




Figure 1: `chatbot_llm` node-level view. The node owns semantic turn interpretation and planner ingress metadata, but it does not create executable plans or dispatch robot skills.




Figure 2: `planner_llm` node-level view. The planner generates structured plan envelopes and maintains feedback-driven goal state; execution and speech remain downstream responsibilities.




Figure 3: `nao_orchestrator` node-level view. Candidate planner admission, direct intents, planner plans, execution feedback, and planner communication remain distinct APIs. Every robot or simulated effect crosses deterministic validation and a registered skill mapping.




Figure 4: Combined people, object-perception, and symbolic-grounding pipeline. ROS4HRI person tracking and object grounding remain separate evidence owners; both contribute to the bounded `grounded_context` consumed by the LLM nodes.